

```
import os
import time
import numpy as np
import random
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras import layers
from sklearn.metrics import confusion_matrix, classification_report
from matplotlib import pyplot as plt
```

```
model_directory = 'models'
```


```
class RandomIntegers:
    def __init__(self):
        pass

    def generate(self, n, length):
        random_integers = random.sample(range(length), n)
        return random_integers
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
def display_images():
    random_integers = RandomIntegers().generate(9, len(x_train))
    plt.figure(figsize=(6, 8))
    counter = 0
    for i in random_integers:
        plt.subplot(330 + 1 + counter)
        counter += 1
        plt.imshow(x_train[i], cmap='gray')
        plt.title(str(y_train[i]))
    plt.show()
```

```
display_images()
```




```
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)
```

```

model = Sequential([
    layers.Conv2D(32, (3, 3), activation = 'relu', input_shape = (28,28,1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

```

 /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape` / `input\_dim` argument to the constructor. Use the `input\_shape` argument of the `compile` method instead.  
super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

```


model.compile(
    optimizer=tf.keras.optimizers.Adam(0.001),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=[tf.keras.metrics.SparseCategoricalAccuracy()]
)

```

```


start_time = time.time()
history = model.fit(
    x_train,
    y_train,
    epochs=25,
    batch_size=512,
    validation_data=(x_test, y_test)
)
end_time = time.time()
time_taken = end_time - start_time
print("Time taken for training: ", time_taken, "seconds")

```

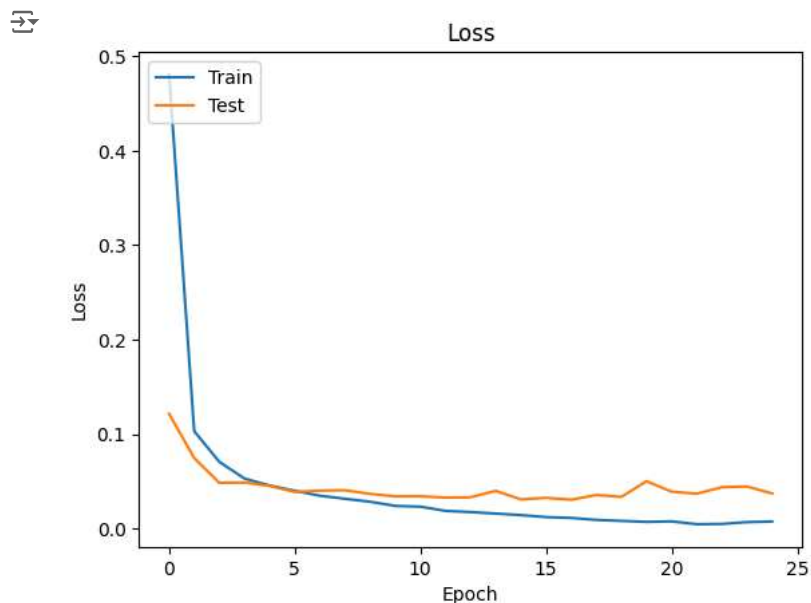
 Epoch 1/25  
/usr/local/lib/python3.10/dist-packages/keras/src/backend/tensorflow/nn.py:609: UserWarning: ``sparse\_categorical\_crossentropy`` receives output, from\_logits = \_get\_logits(  
Epoch 2/25  
118/118 — 52s 432ms/step - loss: 0.9703 - sparse\_categorical\_accuracy: 0.7368 - val\_loss: 0.1216 - val\_sparse\_categorical\_accuracy: 0.7368  
Epoch 3/25  
118/118 — 82s 435ms/step - loss: 0.1141 - sparse\_categorical\_accuracy: 0.9656 - val\_loss: 0.0749 - val\_sparse\_categorical\_accuracy: 0.9656  
Epoch 4/25  
118/118 — 49s 419ms/step - loss: 0.0780 - sparse\_categorical\_accuracy: 0.9749 - val\_loss: 0.0486 - val\_sparse\_categorical\_accuracy: 0.9749  
Epoch 5/25  
118/118 — 82s 421ms/step - loss: 0.0550 - sparse\_categorical\_accuracy: 0.9833 - val\_loss: 0.0490 - val\_sparse\_categorical\_accuracy: 0.9833  
Epoch 6/25  
118/118 — 83s 429ms/step - loss: 0.0488 - sparse\_categorical\_accuracy: 0.9850 - val\_loss: 0.0453 - val\_sparse\_categorical\_accuracy: 0.9850  
Epoch 7/25  
118/118 — 82s 429ms/step - loss: 0.0416 - sparse\_categorical\_accuracy: 0.9868 - val\_loss: 0.0390 - val\_sparse\_categorical\_accuracy: 0.9868  
Epoch 8/25  
118/118 — 82s 425ms/step - loss: 0.0354 - sparse\_categorical\_accuracy: 0.9893 - val\_loss: 0.0404 - val\_sparse\_categorical\_accuracy: 0.9893  
Epoch 9/25  
118/118 — 50s 422ms/step - loss: 0.0345 - sparse\_categorical\_accuracy: 0.9896 - val\_loss: 0.0408 - val\_sparse\_categorical\_accuracy: 0.9896  
Epoch 10/25  
118/118 — 84s 437ms/step - loss: 0.0278 - sparse\_categorical\_accuracy: 0.9914 - val\_loss: 0.0369 - val\_sparse\_categorical\_accuracy: 0.9914  
Epoch 11/25  
118/118 — 83s 447ms/step - loss: 0.0251 - sparse\_categorical\_accuracy: 0.9919 - val\_loss: 0.0343 - val\_sparse\_categorical\_accuracy: 0.9919  
Epoch 12/25  
118/118 — 49s 414ms/step - loss: 0.0231 - sparse\_categorical\_accuracy: 0.9928 - val\_loss: 0.0345 - val\_sparse\_categorical\_accuracy: 0.9928  
Epoch 13/25  
118/118 — 51s 432ms/step - loss: 0.0192 - sparse\_categorical\_accuracy: 0.9937 - val\_loss: 0.0330 - val\_sparse\_categorical\_accuracy: 0.9937  
Epoch 14/25  
118/118 — 82s 431ms/step - loss: 0.0183 - sparse\_categorical\_accuracy: 0.9942 - val\_loss: 0.0334 - val\_sparse\_categorical\_accuracy: 0.9942  
Epoch 15/25  
118/118 — 50s 423ms/step - loss: 0.0168 - sparse\_categorical\_accuracy: 0.9948 - val\_loss: 0.0401 - val\_sparse\_categorical\_accuracy: 0.9948  
Epoch 16/25  
118/118 — 83s 433ms/step - loss: 0.0151 - sparse\_categorical\_accuracy: 0.9953 - val\_loss: 0.0311 - val\_sparse\_categorical\_accuracy: 0.9953  
Epoch 17/25  
118/118 — 82s 432ms/step - loss: 0.0117 - sparse\_categorical\_accuracy: 0.9965 - val\_loss: 0.0328 - val\_sparse\_categorical\_accuracy: 0.9965  
Epoch 18/25  
118/118 — 82s 431ms/step - loss: 0.0104 - sparse\_categorical\_accuracy: 0.9971 - val\_loss: 0.0308 - val\_sparse\_categorical\_accuracy: 0.9971  
Epoch 19/25  
118/118 — 82s 432ms/step - loss: 0.0088 - sparse\_categorical\_accuracy: 0.9975 - val\_loss: 0.0357 - val\_sparse\_categorical\_accuracy: 0.9975  
Epoch 20/25  
118/118 — 82s 434ms/step - loss: 0.0078 - sparse\_categorical\_accuracy: 0.9975 - val\_loss: 0.0338 - val\_sparse\_categorical\_accuracy: 0.9975  
Epoch 21/25  
118/118 — 49s 418ms/step - loss: 0.0062 - sparse\_categorical\_accuracy: 0.9985 - val\_loss: 0.0504 - val\_sparse\_categorical\_accuracy: 0.9985  
Epoch 22/25  
118/118 — 86s 453ms/step - loss: 0.0092 - sparse\_categorical\_accuracy: 0.9973 - val\_loss: 0.0393 - val\_sparse\_categorical\_accuracy: 0.9973  
Epoch 23/25  
118/118 — 79s 428ms/step - loss: 0.0051 - sparse\_categorical\_accuracy: 0.9985 - val\_loss: 0.0371 - val\_sparse\_categorical\_accuracy: 0.9985

```
118/118 ————— 83s 433ms/step - loss: 0.0040 - sparse_categorical_accuracy: 0.9988 - val_loss: 0.0441 - val_sparse_categor
Epoch 24/25
118/118 ————— 81s 425ms/step - loss: 0.0062 - sparse_categorical_accuracy: 0.9980 - val_loss: 0.0447 - val_sparse_categor
Epoch 25/25
118/118 ————— 51s 431ms/step - loss: 0.0068 - sparse_categorical_accuracy: 0.9978 - val_loss: 0.0374 - val_sparse_categor
Time taken for training: 1801.0017359256744 seconds
```

```
plt.plot(history.history['sparse_categorical_accuracy'])
plt.plot(history.history['val_sparse_categorical_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```



```

if not os.path.exists(model_directory):
    os.makedirs(model_directory)
model_path = os.path.join(model_directory, 'model_mnist_cnn.h5')
model.save(model_path)
model.summary()

true_labels = y_test
predicted_labels = np.argmax(model.predict(x_test), axis=-1)

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. Please consider using the SavedModel format instead. For more information, see: https://www.tensorflow.org/guide/saved\_model

Model: "sequential\_1"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d_2 (Conv2D)              | (None, 26, 26, 32) | 320     |
| max_pooling2d_2 (MaxPooling2D) | (None, 13, 13, 32) | 0       |
| conv2d_3 (Conv2D)              | (None, 11, 11, 64) | 18,496  |
| max_pooling2d_3 (MaxPooling2D) | (None, 5, 5, 64)   | 0       |
| flatten_1 (Flatten)            | (None, 1600)       | 0       |
| dense_2 (Dense)                | (None, 64)         | 102,464 |
| dense_3 (Dense)                | (None, 10)         | 650     |


Total params: 365,792 (1.40 MB)  
 Trainable params: 121,930 (476.29 KB)  
 Non-trainable params: 0 (0.00 B)  
 Optimizer params: 243,862 (952.59 KB)  
 313/313 ————— 4s 13ms/step

```
cm = confusion_matrix(true_labels, predicted_labels)
```

```

plt.imshow(cm, cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.xticks(np.arange(10))
plt.yticks(np.arange(10))
plt.show()


```



```

# Generate and print classification report
report = classification_report(true_labels, predicted_labels)
print(report)

```

|  | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
| 0  | 0.99      | 0.99   | 0.99     | 980     |
| 1  | 1.00      | 1.00   | 1.00     | 1135    |
| 2  | 0.99      | 0.99   | 0.99     | 1032    |
| 3  | 0.99      | 1.00   | 0.99     | 1010    |
| 4  | 0.99      | 0.99   | 0.99     | 982     |
| 5  | 0.99      | 0.99   | 0.99     | 892     |
| 6  | 1.00      | 0.98   | 0.99     | 958     |
| 7  | 0.99      | 1.00   | 0.99     | 1028    |
| 8  | 0.99      | 0.99   | 0.99     | 974     |
| 9  | 0.99      | 0.99   | 0.99     | 1009    |
| accuracy   |           |        | 0.99     | 10000   |
| macro avg  | 0.99      | 0.99   | 0.99     | 10000   |
| weighted avg   | 0.99      | 0.99   | 0.99     | 10000   |